a Functional benchmarking

Five functional relations

R1 Add support
R2 Withdraw support
R3 Recompare candidates
R4 Infer consequences
R5 Maintain judgment

→

Benchmark v5

11 sources
6,104 comparison units
13,000 target judgments

→

System diagnosis

7 configurations · 6 families
2 backbones
R1–R3 head transfer

Functional diagnosis motivates operation specialization

b CF-MoA: select the first applicable operation

Current question
Allowed evidence
Initial answer

→

Capability router

first match

- **A3 Candidate comparison** — diagnosis format
- **A1 + A2 Support kernel** — rules apply: add + remove
- **A4 Consequence execution** — supplied model + state
- **A5 Stable task readout (F)** — supported native task
- **Preserve initial output** — otherwise

→

Native output

Sparse execution: one applicable branch per input